Does looking twice help?

We compared two ways to answer questions about the same bookkeeping records. Both use Gemini 3.5 Flash-Lite. The difference is **when the AI gets to see the data.**

Plan up front

Question → Make full plan → Read + calculate → Answer

Inspect, then plan

Question → Read first → Choose next step → Read / calculate → Answer

The first route plans its reads before seeing the results. The second can look, then choose another read or calculation. Server code checks the plan and does the money math; neither route can post transactions through these chat tools. Inspection is bounded: two read rounds, four queries and three model attempts, including repair.

New test questions: frozen automated scores

16 questions per route. **Held-out** means their model answers were not used to adjust the grader before this run.

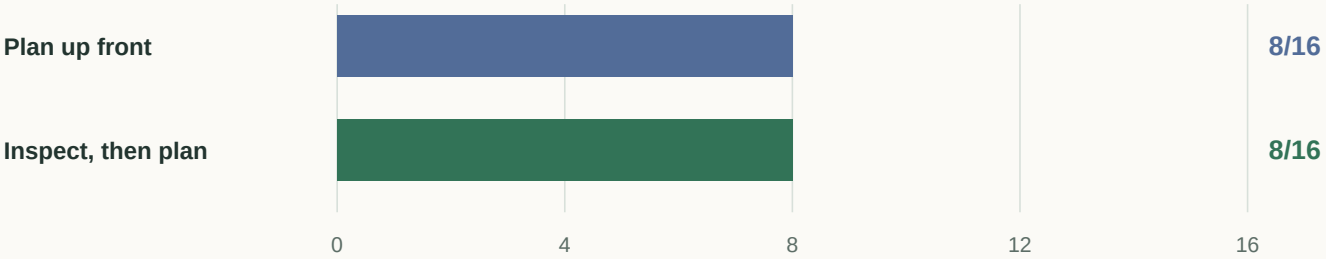

8 questions passed on both routes. 0 passed only after inspection; 0 passed only when planning up front. 8 passed on neither. Every question counts once per route.

Answerable money questions: 4/12 versus 4/12. **Appropriate ask/refuse responses:** 4/4 versus 4/4. The main bars include both; they are not the accuracy rate for numerical answers alone.

Internal review found 4 valid answers rejected by the grader. See the separate review on page 2.

Question type	Plan up front	Inspect first
Direct money questions	4 / 4	4 / 4
Find someone, then compare	0 / 4	0 / 4
Calculations and grouped data	0 / 4	0 / 4
Ask or refuse when appropriate	4 / 4	4 / 4

A pass means the **complete requested result**: exact amounts, the right accounts, dates, currency and business/person, and any requested calculation. The four boundary questions instead require the appropriate clarification or refusal. A plausible paragraph alone does not pass.

Fair comparison, limited conclusion

The same 16 questions ran once through each route, in alternating order, on one isolated synthetic ledger. Expected answers were withheld from the model. The complete configurations differ in tools, instructions and allowed calls, so this does not isolate one causal ingredient.

Plan up front: 19 provider calls; median complete-service time 2.50s. Inspect, then plan: 24 provider calls; median complete-service time 1.92s. Times include all saved executions, including failures; provider calls are not extra independent test questions.

Right numbers are only part of it.

A held-out example from the separate review

Ask: *"Which customer had the most June revenue, and how did that customer's cash movement change from May to June?"* The inspected route found Ledger Labs from **$3,300 revenue**, then read May cash **$900** and June cash **$1,600**: **$1,600 - $900 = +$700**.

An invoice can create revenue before the customer pays: revenue is not cash. The follow-up kept the exact discovered customer identity. The control showed true tables but omitted the change. Internal review judged the inspected answer complete; the frozen grader rejected its table arrangement. That disagreement stays visible below.

Keep grading versions separate

Same held-out answers	Plan up front	Inspect first
Primary: frozen v2 checks	8 / 16	8 / 16
Sensitivity: original v1 checks	8 / 16	8 / 16
Separate AI semantic review	9 / 16	11 / 16

V2 was fixed before held-out calls. Development exposed a grading bug: the control correctly added $675 of supplies and $80 of marketing directly in SQL, but v1 expected a separate calculator step. Corrected development was **3/4 versus 2/4**; original v1 was 2/4 each. Those four questions are not pooled into the 16 held-out scores. Any later AI review stays secondary.

The later internal AI review judged **4 additional responses complete**. The review traced these misses to checks that misread the arrangement of ledger and calculation tables, not wrong arithmetic. 2 of these adaptive successes used one model attempt; do not attribute every gain to extra inspection. The frozen primary score remains unchanged.

Spend and safeguards

Held-out known-use estimate: **$0.2003728**; retained unknown reservation: $0. Add the separate development $0.0430523 / 10 calls once: **$0.2434251 across 53 calls**. Caps stayed $0.40 / 80 for holdout and $0.50 overall. These are tracked standard-rate estimates, not an invoice or account balance.

The offline integrity audit passed 866 checks: fixed source, all 32 outcomes, retained errors, usage and before/after accounting fingerprints. The test ledger stayed unchanged; this says nothing about the current hosted books.

What would justify a next release check?

The predefined screen requires at least two net extra held-out successes, no lost common-question success, and no additional wrong financial answer or hard-boundary violation. Passing that screen only permits separate hosted acceptance; it does not automatically change the deployed default.

Decision: retain the current plan-up-front default. The adaptive route did not clear the required two-success gain on the frozen primary score. A secondary review does not retroactively change that screen.

Sixteen purpose-built cases are a small sample. The questions and reference answers were authored with AI; semantic review is internal AI review, not independent human accounting approval. Structured grading is not a proof of every sentence. One run cannot establish universal accuracy, stable provider superiority, browser quality or million-user capacity.

What still broke

Some answers skipped the requested calculation. Two adaptive answers copied complicated row references incorrectly. Two requests used an invalid operand field even after one schema-repair attempt. Those failures were preserved. The next proposed repair is simpler row references and more specific validation feedback, while keeping strict money checks. These repairs are not implemented or evaluated by this report.

Evidence: run 41258129206; commit 84eb5de791b5; evaluator 1b6f4642c68019a8...; primary adaptive-scoring-v2. Full bindings accompany the PDF in its render receipt.