Can it learn how you keep your books?

Think of a category as a labelled drawer. Paper might go in **Office supplies**. Your business might prefer a separate **Printing expense** drawer. We tested whether three different helpers can understand that preference without overusing it.

One example, three moments

Moment	What happens	What it should mean
1. Buy paper	The helper suggests **Office supplies**.	That can be a reasonable answer before you teach your preference.
2. Review + post	You change the suggestion to **Printing expense**, review the entry, then post it.	A trusted example is saved. A later similar paper purchase can use it. A draft alone does not teach.
3. Buy a laptop	The same shop sells you a laptop.	It must not inherit the paper category. Equipment needs full entry review; this small selector hands it off.

This is an illustration of intended behavior, not a claim that any tested helper passed it.

The three helpers got the same information

Helper	How it chooses
Exact-memory rules	Ordinary code checks an exact description + seller match, then simple account-label rules. No paid model.
Gemini	A language model reads the facts and examples, then chooses an allowed category, asks a question, or hands off.
Jev	A decision model chooses among the same allowed categories, clarification choices and handoffs. Its confidence is not measured accuracy.

Fixed facts
Amount + payment account

>

Suggest or ask
Same accounts + examples

>

You review
Only posting teaches

What stays under our control: code fixes the amount, date, currency, rate and payment account; the server supplies the allowed accounts and eligible history. The helper chooses a category. It cannot post money by itself.

Learning here means using reviewed examples. We are not retraining either model. The test gives each helper the same transaction after **0, 1 and 3 earlier review events**. Old, reversed or another business's examples can be excluded, so three events need not mean three usable examples.

The narrow selector covers simple expense and revenue categories. Invoices, loans, refunds, transfers, assets and complex entries need the separate full-entry/manual workflow.

Two setup mistakes stay in the record. Before held-out testing, our adapter rejected extra Gemini billing information in v1 and rounded Jev confidence in v2. We fixed these integration checks and kept both failures. They are not evidence that the models chose wrong categories, and their scores are not mixed into v3.

Retained development costs: **v1** Gemini 1 attempt / $0.000257100000; **v2** Gemini 2 attempts / $0.000838300000 and Jev 2 attempts / $0.000188076000. These stay inside the same total allowance.

What did the test show?

18 held-out stories × 3 review states. Each helper faced 54 inputs. There were 33 inputs needing a category, covering 11 distinct stories; errors stay in the denominator.

1 / Accuracy: did it choose a sensible category?

Counts above accept any pre-labelled reasonable category. Following a taught business preference is a separate, stricter check below. These are category scores, not whole-journal accuracy.

2 / Learning: did it follow the taught preference?

Helper	Before reviews	After 3 reviews	Category + preference
Exact-memory rules	0/5	5/5	16/33
Gemini	2/5	5/5	33/33
Jev	2/5	4/5	31/33

Before/after uses the same **five teaching stories**. A different reasonable category before teaching is not an error. The final column is the stricter score across all 33 category inputs.

A real miss: Jev kept **Office supplies** after both one and three reviewed paper purchases, instead of the taught **Printing** preference. The category was reasonable; the preference was missed twice.

3 / Questions, failures and cost

Helper	Needed questions	Unneeded questions	Run errors	Estimated cost + unknown reserve
Exact-memory rules	9/9	17	0	$0 + $0 held
Gemini	9/9	0	0	$0.021673 + $0 held
Jev	9/9	0	0	$0.004771 + $0 held

Questions include shared code checks. Run errors mean execution failed, not a wrong category. Cost is this held-out phase only: usage estimates + unknown-use reserves, not an invoice. Development and separate full-entry acceptance are excluded.

V3 development + held-out: Gemini $0.02904; Jev $0.00643. **All three versions:** Gemini $0.030135 / 60 attempts; Jev $0.006618 / 59 attempts. Totals include $0 in unknown-use reserves; each earlier cost is counted once.

What should we choose?

Keep Gemini for this narrow selector for now. This small test did not meet the pre-set rule for replacing it with Jev. That does not prove Gemini is universally better.

Jev minus Gemini on category + preference: **-6.1 percentage points**; paired 95% interval **-18.2 to +0.0**, across 11 story groups. Jev's tracked cost per correct category + preference was 76.6% lower. The pre-set screen still requires enough evidence against a meaningful quality loss.

Rules / Gemini / Jev: unsafe suggestions **0 / 0 / 0**; harmful carry-over **0 / 0 / 0**. Zero observed failures would not prove zero risk.

Keep the claim small: these are synthetic, AI-authored stories with no independent human adjudication. The three review states are related examples, not 54 independent users. Intervals are exploratory and unadjusted. Context was supplied equally; database retrieval, posting safety and full-entry accuracy are separate tests.