# Did it answer the whole question?

We gave the same Gemini model the same invented bookkeeping records. Sixteen unseen questions were each asked three ways: 48 service runs. “Unseen” means these questions were not used to adjust the candidate or frozen primary checker before the test. They were authored with AI, not sampled from real customers.

| Three ways to answer | What the assistant does |
|---|---|
| Plan up front | Decides what to read and calculate before seeing the amounts. |
| Inspect then plan | Can read the amounts, then decide what to read or calculate next. |
| Track requested results | Also lists requested results; code checks the list and can finish some declared calculations. |

## New questions: frozen checker scores

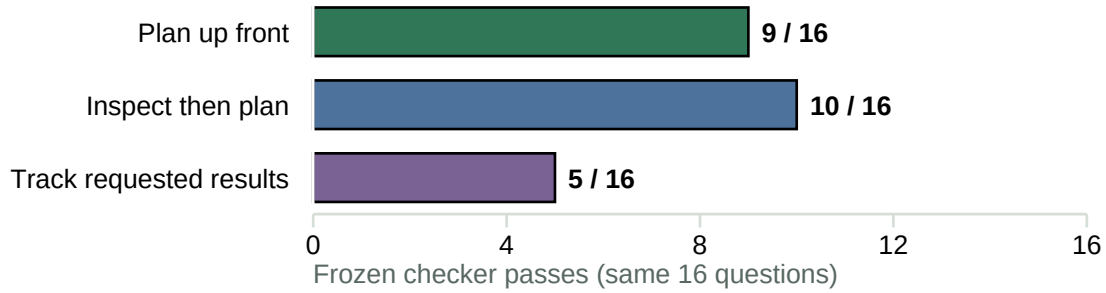


Four questions require asking for information or refusing an unsupported request. Those count in the 16. On the 12 answerable money questions: Plan up front 6/12; Inspect then plan 7/12; Track requested results 1/12. The checker has documented grading errors; see the separate check below.

| Question type | Plan up front | Inspect then plan | Track requested results |
|---|---|---|---|
| Direct money questions | 4 / 4 | 4 / 4 | 1 / 4 |
| Find someone, then compare | 2 / 4 | 2 / 4 | 0 / 4 |
| Calculations and grouped data | 0 / 4 | 1 / 4 | 0 / 4 |
| Ask or refuse when needed | 3 / 4 | 3 / 4 | 4 / 4 |

## How the checker awards a pass

A separate checker uses expected numbers and the actual book rows. It checks accounts, dates, currency, source, requested calculations and missing parts. It also checks that the main answer focused on the requested question: a correct total for all expenses does not answer “office supplies.” The assistant’s own “complete” label earns no credit.

## Separate check: same answers, two grader fixes

| Plan up front | Inspect then plan | Track requested results |
|---|---|---|
| 11 / 16 | 11 / 16 | 5 / 16 |

Fixing how the checker reads monthly units and recorded user assumptions changes these counts. The model answers were not rerun. This later diagnostic is separate from the frozen scores and does not change the decision.

# What this result means

## Keep Plan up front.

The live app at report time uses **Plan up front** (source 6691d76ad245). A comparison does not itself deploy a change.

The decision rule was set before testing. A candidate needs at least two more complete answers, no detected wrong-money or safety failures, no lost direct-question success and all four boundary cases correct. A paired statistical test must also pass, adjusted for checking two candidates. If that screen is inconclusive, keep Plan up front.

| Compared with Plan up front | Net extra passes | Adjusted p | Screen |
|---|---|---|---|
| Inspect then plan | +1 | 1.000 | Not qualified |
| Track requested results | -4 | 1.000 | Not qualified |

Adjusted p is a check against mistaking chance for an improvement; the required cutoff is 0.05. It is not the probability that an answer is right.

## One actual question, traced to the records

*“Please give me posted Site operations expenses for each month from July through September 2026, including any month with no activity.”*

The answer showed July $105, August $0 and September $0. The posted Site operations lines are $55 + $35 + $15 in July, with no scoped postings in the other two months. It included archived subaccount history and displayed all three requested months. This is one checked example, not evidence that every answer is correct.

Example H01, Plan up front; frozen score: pass. Explanation checked by an internal AI reviewer, not a human accountant.

## Separate check on the hosted app

Three known questions checked Office supplies over all history, October supplies, and Rent minus supplies. The original automatic checks passed 2/3; internal AI review verified 3/3. The remaining automatic failure was a redundant USD label. Original results were preserved. These known cases are not a new accuracy score.

## Costs and limits

This held-out run: **$0.2652893 known-use estimate + $0 reserved for unknown usage**, across 52 model calls. The tracked allowance, including development and preserved failures: **$0.6919872** ($0.3796494 known + $0.3123378 unknown), within $0.90. These are tracked estimates/reservations, not an invoice.

The known development smoke scored 0/3 and remains separate, with its costs preserved. Routes share one unchanged ledger but differ in tools, instructions and allowed calls; this compares whole setups, not one cause. AI wrote the cases and references; no independent human accountant judged them. Related questions may be correlated. This small test cannot establish universal accuracy or production-scale reliability.

Evidence: run 41282205719; tested source 6691d76ad245; frozen completeness grader. Full source hashes and review bindings accompany this PDF.